



# **Introduction to Data Science Project**

## **Predict Student's dropout and academic success**

Project report submitted By  
**Group-9**

Tushar Sharma – 21UCC104

Rajdeep Tiwari – 21UCS166

Sahil Sidana – 21UCS177

Shreyansh Jain – 21UCS198

Course Instructors

Dr. Lal Upendra Pratap Singh

Dr. Subrat Kumar Dash

Dr. Aloke Datta

# Contents

1. Objective .....	3
2. Dataset Description .....	3
2.1. About Data .....	3
2.2. About rows and features .....	5
3. Imported Libraries .....	6
4. Data Preprocessing .....	7
5. Exploratory Data Analysis .....	10
Categorical Data.....	10
5.1. Target .....	10
5.2. Gender .....	11
5.3. Marital Status .....	12
5.4. Debtor .....	13
5.5. Scholarship holder .....	14
5.6. Tuition Fee up to Date .....	15
5.7. Age at Enrollment .....	16
5.8. Nationality .....	18
5.9. Application Mode .....	19
5.10. Daytime/Evening Attendance .....	21
5.11. Course .....	22
5.12. Previous Qualifications .....	24
Continuous data.....	25
6. Feature Engineering.....	27
7. Model Building .....	30
8. Models .....	31
8.1. Naive Bayes .....	31
8.2. Logistic Regression .....	33
8.3. Decision Tree .....	35
8.4. Random Forest .....	37
8.5. Support Vector Machine (SVM) .....	39
9. Model Evaluation and Conclusion .....	41
9.1. Result Matrix .....	41
9.2. ROC Graph .....	42
9.3. Conclusion .....	42
10. References.....	43

**Google colab Link:** [IDS Project](#)

## **1. Objective**

Our problem statement is to use the given Academic Background of a student and to predict the academic dropout and failure in higher education, by using machine learning techniques to identify students at risk at an early stage of their academic path, so that strategies to support them can be put into place.

## **2. Dataset Description**

**Source of Dataset:** [Predict students' dropout and academic success](#)

A dataset created from a higher education institution (acquired from several disjoint databases) related to students enrolled in different undergraduate degrees, such as agronomy, design, education, nursing, journalism, management, social service, and technologies. The dataset includes information known at the time of student enrollment (academic path, demographics, and social-economic factors) and the students' academic performance at the end of the first and second semesters. The data is used to build classification models to predict students' dropout and academic success. The problem is formulated as a three category classification task (dropout, enrolled, and graduate) at the end of the normal duration of the course.

# Show first five rows.  
data.head()

	Marital Status	Application mode	Application order	Course	Daytime/evening attendance	Previous qualification	Previous qualification (grade)	Nacionality	Mother's qualification	Father's qualification	Mother's occupation	Father's occupation	Admission grade
0	1	17	5	171	1	1	122.0	1	19	12	5	9	1
1	1	15	1	9254	1	1	160.0	1	1	3	3	3	1
2	1	1	5	9070	1	1	122.0	1	37	37	9	9	1
3	1	17	2	9773	1	1	122.0	1	38	37	5	3	1
4	2	39	1	8014	0	1	100.0	1	37	38	9	9	1

[ ] # Show last five rows.  
data.tail()

	Marital Status	Application mode	Application order	Course	Daytime/evening attendance	Previous qualification	Previous qualification (grade)	Nacionality	Mother's qualification	Father's qualification	Mother's occupation	Father's occupation	Admission grade
4419	1	1	6	9773	1	1	125.0	1	1	1	5	4	
4420	1	1	2	9773	1	1	120.0	105	1	1	9	9	
4421	1	1	1	9500	1	1	154.0	1	37	37	9	9	
4422	1	1	1	9147	1	1	180.0	1	37	37	7	4	
4423	1	10	1	9773	1	1	152.0	22	38	37	5	9	

data.describe()

	Marital Status	Application mode	Application order	Course	Daytime/evening attendance	Previous qualification	Previous qualification (grade)	Nacionality	Mother's qualification	Father's qualification	Mother's occupation	Father's occupation
count	4424.000000	4424.000000	4424.000000	4424.000000	4424.000000	4424.000000	4424.000000	4424.000000	4424.000000	4424.000000	4424.000000	4424.000000
mean	1.178571	18.669078	1.727848	8856.642631	0.890823	4.577758	132.613314	1.873192	19.561935	22.275316	10.960895	11.000000
std	0.605747	17.484682	1.313793	2063.566416	0.311897	10.216592	13.188332	6.914514	15.603186	15.343108	26.418253	25.000000
min	1.000000	1.000000	0.000000	33.000000	0.000000	1.000000	95.000000	1.000000	1.000000	1.000000	0.000000	0.000000
25%	1.000000	1.000000	1.000000	9085.000000	1.000000	1.000000	125.000000	1.000000	2.000000	3.000000	4.000000	4.000000
50%	1.000000	17.000000	1.000000	9238.000000	1.000000	1.000000	133.100000	1.000000	19.000000	19.000000	5.000000	7.000000
75%	1.000000	39.000000	2.000000	9556.000000	1.000000	1.000000	140.000000	1.000000	37.000000	37.000000	9.000000	9.000000
max	6.000000	57.000000	9.000000	9991.000000	1.000000	43.000000	190.000000	109.000000	44.000000	44.000000	194.000000	195.000000

## It contains 4424 rows and 37 columns/features

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 4424 entries, 0 to 4423
```

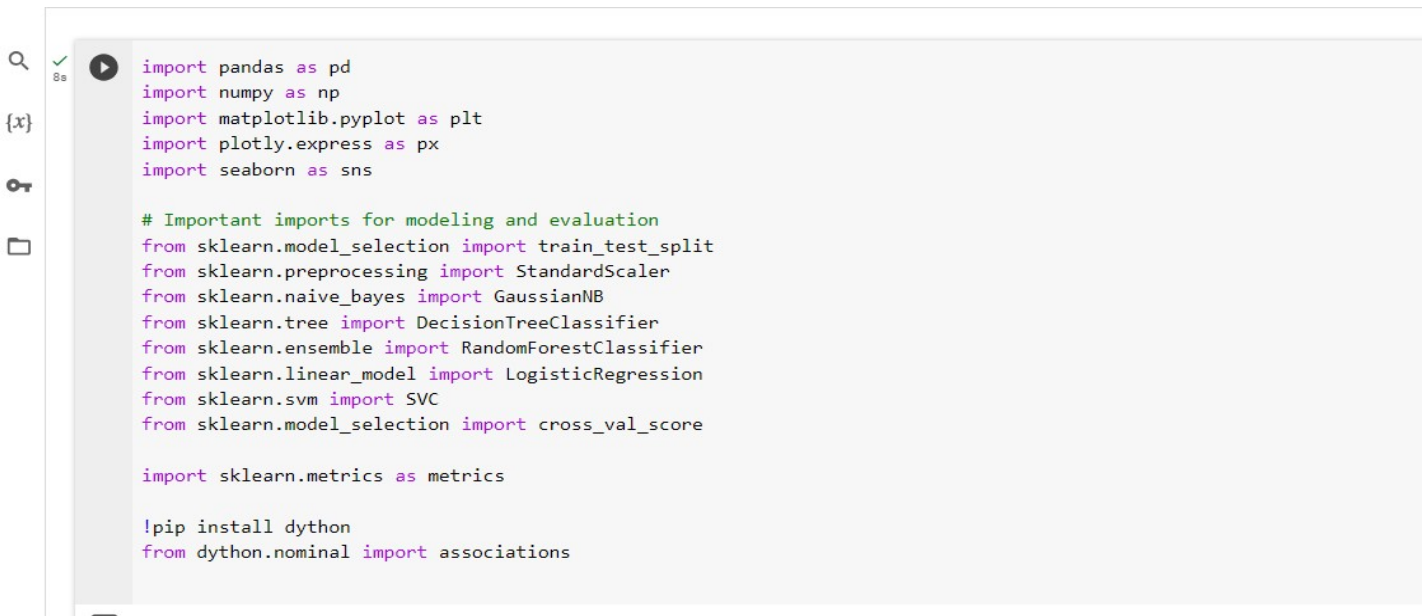
```
Data columns (total 37 columns):
```

#	Column	Non-Null Count	Dtype
0	Marital Status	4424 non-null	int64
1	Application mode	4424 non-null	int64
2	Application order	4424 non-null	int64
3	Course	4424 non-null	int64
4	Daytime/evening attendance	4424 non-null	int64
5	Previous qualification	4424 non-null	int64
6	Previous qualification (grade)	4424 non-null	float64
7	Nacionality	4424 non-null	int64
8	Mother's qualification	4424 non-null	int64
9	Father's qualification	4424 non-null	int64
10	Mother's occupation	4424 non-null	int64
11	Father's occupation	4424 non-null	int64
12	Admission grade	4424 non-null	float64
13	Displaced	4424 non-null	int64
14	Educational special needs	4424 non-null	int64
15	Debtor	4424 non-null	int64
16	Tuition fees up to date	4424 non-null	int64
17	Gender	4424 non-null	int64
18	Scholarship holder	4424 non-null	int64
19	Age at enrollment	4424 non-null	int64
20	International	4424 non-null	int64
21	Curricular units 1st sem (credited)	4424 non-null	int64
22	Curricular units 1st sem (enrolled)	4424 non-null	int64
23	Curricular units 1st sem (evaluations)	4424 non-null	int64
24	Curricular units 1st sem (approved)	4424 non-null	int64
25	Curricular units 1st sem (grade)	4424 non-null	float64
26	Curricular units 1st sem (without evaluations)	4424 non-null	int64
27	Curricular units 2nd sem (credited)	4424 non-null	int64
28	Curricular units 2nd sem (enrolled)	4424 non-null	int64
29	Curricular units 2nd sem (evaluations)	4424 non-null	int64
30	Curricular units 2nd sem (approved)	4424 non-null	int64
31	Curricular units 2nd sem (grade)	4424 non-null	float64
32	Curricular units 2nd sem (without evaluations)	4424 non-null	int64
33	Unemployment rate	4424 non-null	float64
34	Inflation rate	4424 non-null	float64
35	GDP	4424 non-null	float64
36	Target	4424 non-null	object

```
dtypes: float64(7), int64(29), object(1)
```

```
memory usage: 1.2+ MB
```

## 3. Imported Libraries

A screenshot of a code editor interface. On the left, there is a sidebar with icons for search, a file explorer, and a key symbol. The main area displays Python code with syntax highlighting. The code imports several libraries: pandas as pd, numpy as np, matplotlib.pyplot as plt, plotly.express as px, and seaborn as sns. It then lists important imports for modeling and evaluation, including train\_test\_split, StandardScaler, GaussianNB, DecisionTreeClassifier, RandomForestClassifier, LogisticRegression, SVC, and cross\_val\_score from sklearn. It also imports metrics from sklearn.metrics. At the bottom, there is a command to install dython and an import for associations from dython.nominal.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import plotly.express as px
import seaborn as sns

# Important imports for modeling and evaluation
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.model_selection import cross_val_score

import sklearn.metrics as metrics

!pip install dython
from dython.nominal import associations
```

**3.1 pandas:** Data manipulation and analysis library in Python, providing easy-to-use data structures like DataFrames for working with structured data.

**3.2 numpy:** Numerical computing library in Python, offering support for large, multi-dimensional arrays and matrices, along with mathematical functions to operate on them.

**3.3 seaborn:** Statistical data visualization library built on top of matplotlib, providing an aesthetically pleasing interface for creating informative and attractive statistical graphics.

**3.4 scikit-learn (sklearn):** Machine learning library in Python, offering simple and efficient tools for data analysis and modeling, including classification, regression, clustering, and dimensionality reduction algorithms.

**3.5 dython:** A library for data analysis and feature engineering, specifically designed to work seamlessly with pandas, providing additional tools for handling missing data, encoding categorical features, and exploring data relationships.

**3.6 matplotlib:** 2D plotting library in Python, producing static, animated, and interactive visualizations in a variety of formats, and serving as the foundation for other visualization libraries like seaborn.

## 4. Data Preprocessing:

- Correcting column name from Nacionality to Nationality

✓ [7] #DATA PREPROCESSING.

0s

```
# Rename column Nacionality.
data.rename(columns={'Nacionality': 'Nationality'}, inplace=True)
```

```
data.columns
```

```
Index(['Marital Status', 'Application mode', 'Application order', 'Course',
      'Daytime/evening attendance', 'Previous qualification',
      'Previous qualification (grade)', 'Nationality',
      'Mother's qualification', 'Father's qualification',
      'Mother's occupation', 'Father's occupation', 'Admission grade',
      'Displaced', 'Educational special needs', 'Debtor',
      'Tuition fees up to date', 'Gender', 'Scholarship holder',
      'Age at enrollment', 'International',
      'Curricular units 1st sem (credited)',
      'Curricular units 1st sem (enrolled)',
      'Curricular units 1st sem (evaluations)',
      'Curricular units 1st sem (approved)',
      'Curricular units 1st sem (grade)',
      'Curricular units 1st sem (without evaluations)',
      'Curricular units 2nd sem (credited)',
      'Curricular units 2nd sem (enrolled)',
      'Curricular units 2nd sem (evaluations)',
      'Curricular units 2nd sem (approved)',
      'Curricular units 2nd sem (grade)',
      'Curricular units 2nd sem (without evaluations)', 'Unemployment rate',
      'Inflation rate', 'GDP', 'Target'],
      dtype='object')
```

- Changing categorical columns to category datatype.

```

0s cont_cols = ['Age at enrollment', 'Curricular units 1st sem (credited)', 'Curricular units 1st sem (enrolled)',
               'Curricular units 1st sem (evaluations)', 'Curricular units 1st sem (approved)', 'Curricular units 1st sem (grade)',
               'Curricular units 1st sem (without evaluations)', 'Curricular units 2nd sem (credited)',
               'Curricular units 2nd sem (enrolled)', 'Curricular units 2nd sem (evaluations)',
               'Curricular units 2nd sem (approved)', 'Curricular units 2nd sem (grade)',
               'Curricular units 2nd sem (without evaluations)', 'Unemployment rate', 'Inflation rate', 'GDP', 'Target']

# Get all categorical variables except target.
cat_cols = ['Marital Status', 'Application mode', 'Application order', 'Course', 'Daytime/evening attendance',
            'Previous qualification', 'Nationality', 'Mother's qualification', 'Father's qualification',
            'Mother's occupation', 'Father's occupation', 'Displaced', 'Educational special needs', 'Debtor',
            'Tuition fees up to date', 'Gender', 'Scholarship holder', 'International']

# Change categorical columns to category datatype.
data[cat_cols] = data[cat_cols].astype('category')

data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4424 entries, 0 to 4423
Data columns (total 37 columns):
 #   Column                                     Non-Null Count  Dtype
---  -
 0   Marital Status                           4424 non-null   category
 1   Application mode                          4424 non-null   category
 2   Application order                        4424 non-null   category
 3   Course                                  4424 non-null   category
 4   Daytime/evening attendance               4424 non-null   category
 5   Previous qualification                   4424 non-null   category
 6   Previous qualification (grade)           4424 non-null   float64
 7   Nationality                             4424 non-null   category
 8   Mother's qualification                   4424 non-null   category
 9   Father's qualification                   4424 non-null   category
10  Mother's occupation                       4424 non-null   category
11  Father's occupation                       4424 non-null   category
12  Admission grade                           4424 non-null   float64
13  Displaced                                4424 non-null   category
14  Educational special needs                 4424 non-null   category
15  Debtor                                   4424 non-null   category
16  Tuition fees up to date                   4424 non-null   category
17  Gender                                   4424 non-null   category
18  Scholarship holder                       4424 non-null   category
19  Age at enrollment                        4424 non-null   int64
20  International                             4424 non-null   category
21  Curricular units 1st sem (credited)       4424 non-null   int64
22  Curricular units 1st sem (enrolled)       4424 non-null   int64
23  Curricular units 1st sem (evaluations)    4424 non-null   int64
24  Curricular units 1st sem (approved)       4424 non-null   int64
25  Curricular units 1st sem (grade)          4424 non-null   float64
26  Curricular units 1st sem (without evaluations) 4424 non-null   int64
27  Curricular units 2nd sem (credited)       4424 non-null   int64
28  Curricular units 2nd sem (enrolled)       4424 non-null   int64
29  Curricular units 2nd sem (evaluations)    4424 non-null   int64
30  Curricular units 2nd sem (approved)       4424 non-null   int64
31  Curricular units 2nd sem (grade)          4424 non-null   float64
32  Curricular units 2nd sem (without evaluations) 4424 non-null   int64
33  Unemployment rate                         4424 non-null   float64
34  Inflation rate                           4424 non-null   float64
35  GDP                                       4424 non-null   float64
36  Target                                   4424 non-null   object
dtypes: category(18), float64(7), int64(11), object(1)
memory usage: 744.1+ KB

```



- Check for missing values in all columns

Q	# Check for missing values. data.isna().sum()	
{x}		
key		
file		
	Marital Status	0
	Application mode	0
	Application order	0
	Course	0
	Daytime/evening attendance	0
	Previous qualification	0
	Previous qualification (grade)	0
	Nationality	0
	Mother's qualification	0
	Father's qualification	0
	Mother's occupation	0
	Father's occupation	0
	Admission grade	0
	Displaced	0
	Educational special needs	0
	Debtor	0
	Tuition fees up to date	0
	Gender	0
	Scholarship holder	0
	Age at enrollment	0
	International	0
	Curricular units 1st sem (credited)	0
	Curricular units 1st sem (enrolled)	0
	Curricular units 1st sem (evaluations)	0
	Curricular units 1st sem (approved)	0
	Curricular units 1st sem (grade)	0
	Curricular units 1st sem (without evaluations)	0
	Curricular units 2nd sem (credited)	0
	Curricular units 2nd sem (enrolled)	0
	Curricular units 2nd sem (evaluations)	0
	Curricular units 2nd sem (approved)	0
	Curricular units 2nd sem (grade)	0
	Curricular units 2nd sem (without evaluations)	0
	Unemployment rate	0
	Inflation rate	0
<>	GDP	0
list	Target	0
—	dtype: int64	

Note: There are no missing values

## 5. Exploratory Data Analysis

### For Categorical data

```
[10] #Exploratory data analysis
new_data = data.copy()

# Show all unique values of target variable.
new_data['Target'].unique()

array(['Dropout', 'Graduate', 'Enrolled'], dtype=object)
```

### 5.1 Target

Creating a pie chart that depicts the proportions of total number of dropouts, total number of graduated, and total number of enrolled students.


```
[11] x = new_data['Target'].value_counts().index
y = new_data['Target'].value_counts().values

df = pd.DataFrame({
    'Target': x,
    'Count_T': y
})

fig = px.pie(df,
             names='Target',
             values='Count_T',
             title='How many dropouts, enrolled & graduates are there in Target column')

fig.update_traces(labels=['Graduate', 'Dropout', 'Enrolled'], hole=0.4, textinfo='value+label', pull=[0, 0.2, 0.1])
fig.show()
```

How many dropouts, enrolled & graduates are there in Target column



The number of graduated students is more than the number of dropout students.

The total number of graduated and dropout students is  $2209 + 1421 = 3630$ , which is the number of observations for building our model.

## 5.2 Gender

Creating a pie chart that depicts the proportions of total number of female, and total number of male students who dropped out.

```
[12] # Gender

dropout_data = new_data.loc[new_data['Target'] == "Dropout"]

# Count the occurrences of each marital status
gender_counts = dropout_data['Gender'].value_counts().reset_index()

# Rename the columns for clarity in the plot
gender_counts.columns = ['Gender', 'Count']

# Calculate percentages
gender_counts['Percentage'] = (gender_counts['Count'] / gender_counts['Count'].sum()) * 100


# Plotting the pie chart using plotly.express
fig = px.pie(gender_counts, names='Gender', values='Percentage', title='Gender Distribution for Dropout Students', color_discrete_sequence=px.colors.qualitative.Set3)
fig.update_traces(labels=['Male', 'Female'], hole=0.4, textinfo='percent+label', pull=[0.1,0.1,0.1,0.1])

fig.show()

# 0 : Male
# 1 : Female

/usr/local/lib/python3.10/dist-packages/numpy/core/numeric.py:2463: FutureWarning:
elementwise comparison failed; returning scalar instead, but in the future will perform elementwise comparison
```

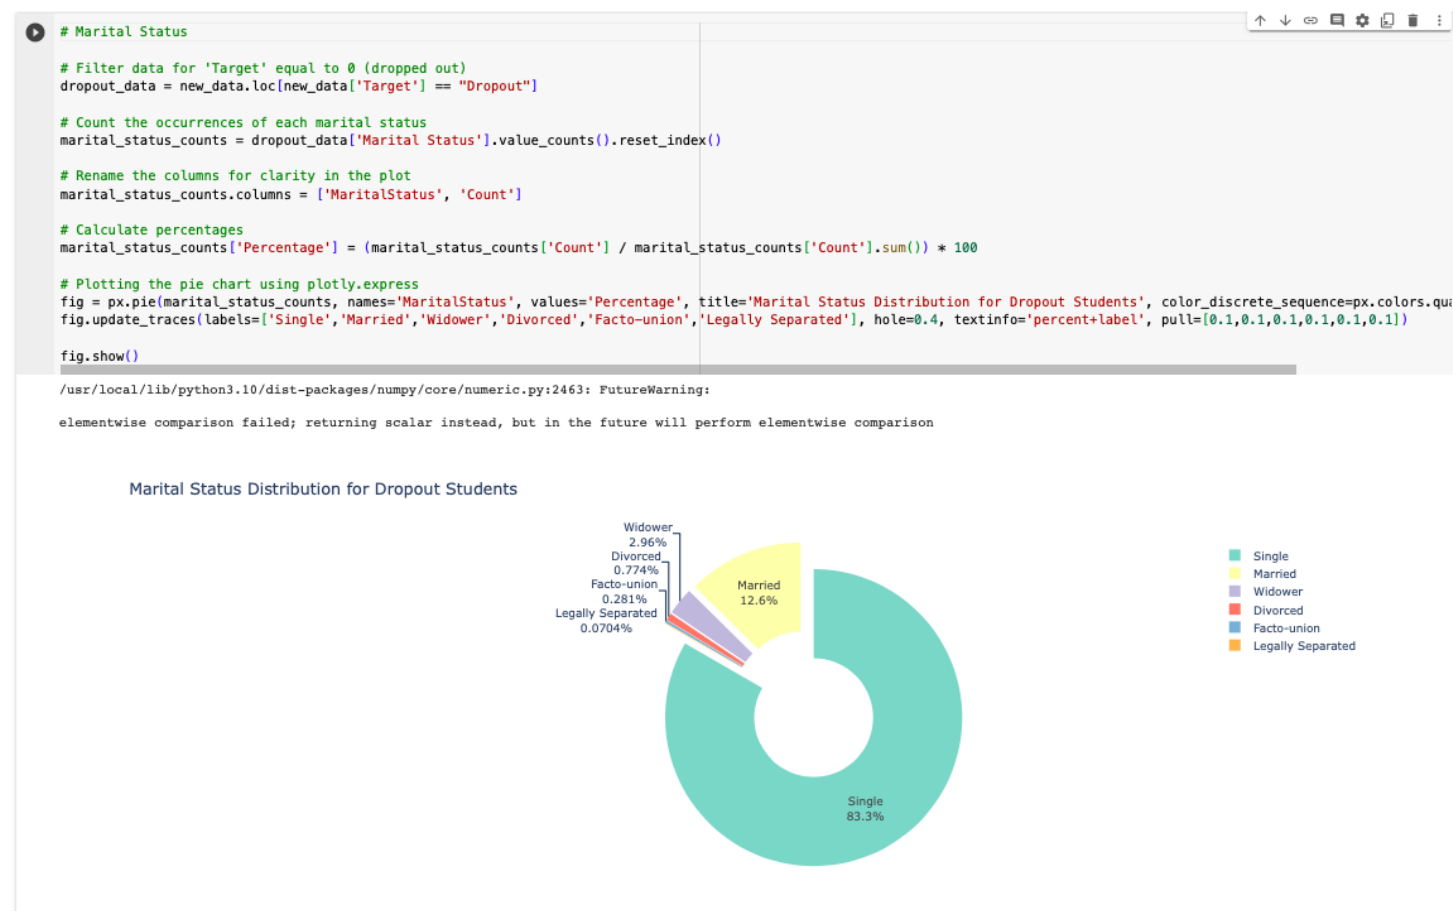
Gender Distribution for Dropout Students



The number of dropout female and male students is almost equal.  
 There are more female students than male students in the dataset.  
 The dataset is slightly imbalanced, however it should not significantly affect the future model.

## 5.3 Marital Status

Creating a pie chart that shows the number of students in the following marital status groups: single, married, divorced, facto union, widower, and legally who dropped out.



Number of Students with a marital status Single who have dropped out of the academic institutions is the highest among the total dropouts .

## 5.4 Debtor

Creating a bar chart for the dropped out students who have a debt or not



## 5.5 Scholarship Holder

Creating a bar chart for the dropped out students who have a scholarship or not

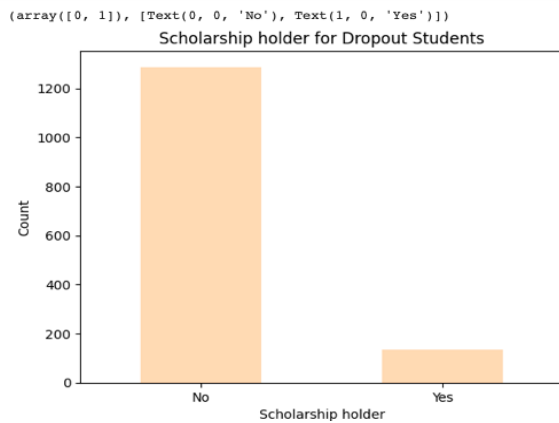
```
✓ [15] # Scholarship Holder
0s
new_data['Scholarship holder']=new_data['Scholarship holder'].replace({
    0: 'No',
    1: 'Yes'
})

filtered_data = new_data[new_data['Target'] == 'Dropout']

sorted_data = filtered_data.sort_values(by='Scholarship holder', ascending=False)

value_counts = sorted_data['Scholarship holder'].value_counts()

value_counts.plot(kind='bar', color='peachpuff')
plt.xlabel('Scholarship holder')
plt.ylabel('Count')
plt.title('Scholarship holder for Dropout Students')
plt.xticks(rotation = 0)
```



## 5.6 Tuition fees up to Date

Creating a bar chart for the dropped out students who have paid tuition fees up to date or not

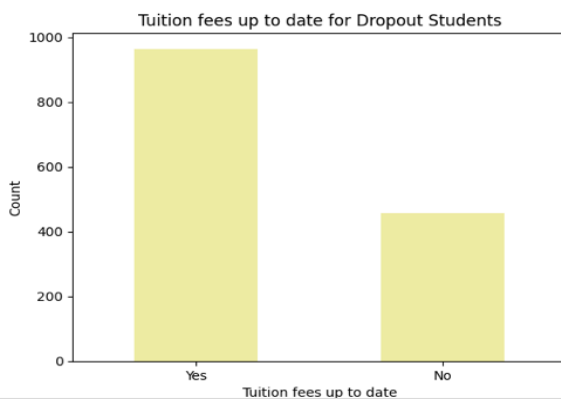
```
[16] # Tuition Fees
new_data['Tuition fees up to date']=new_data['Tuition fees up to date'].replace({
    0: 'No',
    1: 'Yes'
})

filtered_data = new_data[new_data['Target'] == 'Dropout']

sorted_data = filtered_data.sort_values(by='Tuition fees up to date', ascending=False)

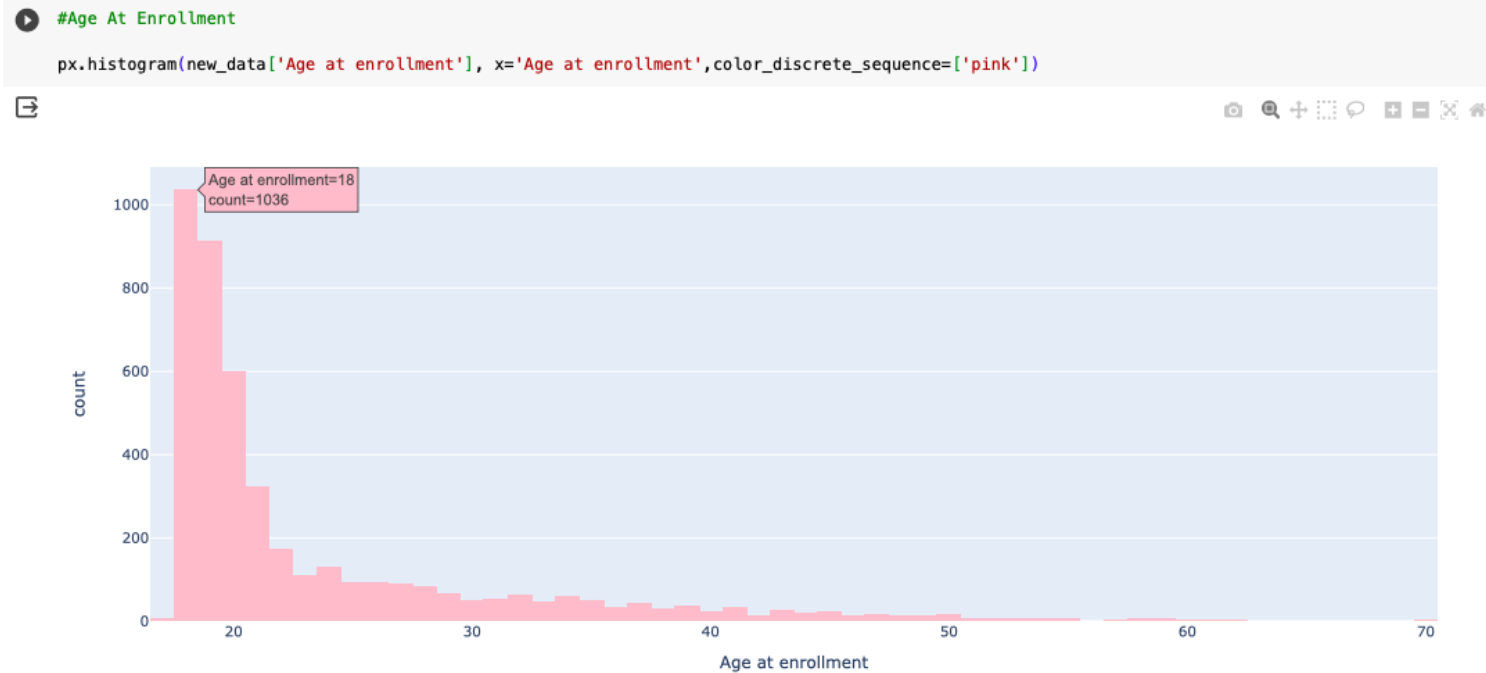
value_counts = sorted_data['Tuition fees up to date'].value_counts()

value_counts.plot(kind='bar', color='palegoldenrod')
plt.xlabel('Tuition fees up to date')
plt.ylabel('Count')
plt.title('Tuition fees up to date for Dropout Students')
plt.xticks(rotation=0)
plt.show()
```



## 5.7 Age at Enrollment

Creating a histogram to depict the age at enrollment distribution



The vast majority of students are 17-22 years old. The number of students decreases as the student's age increases.

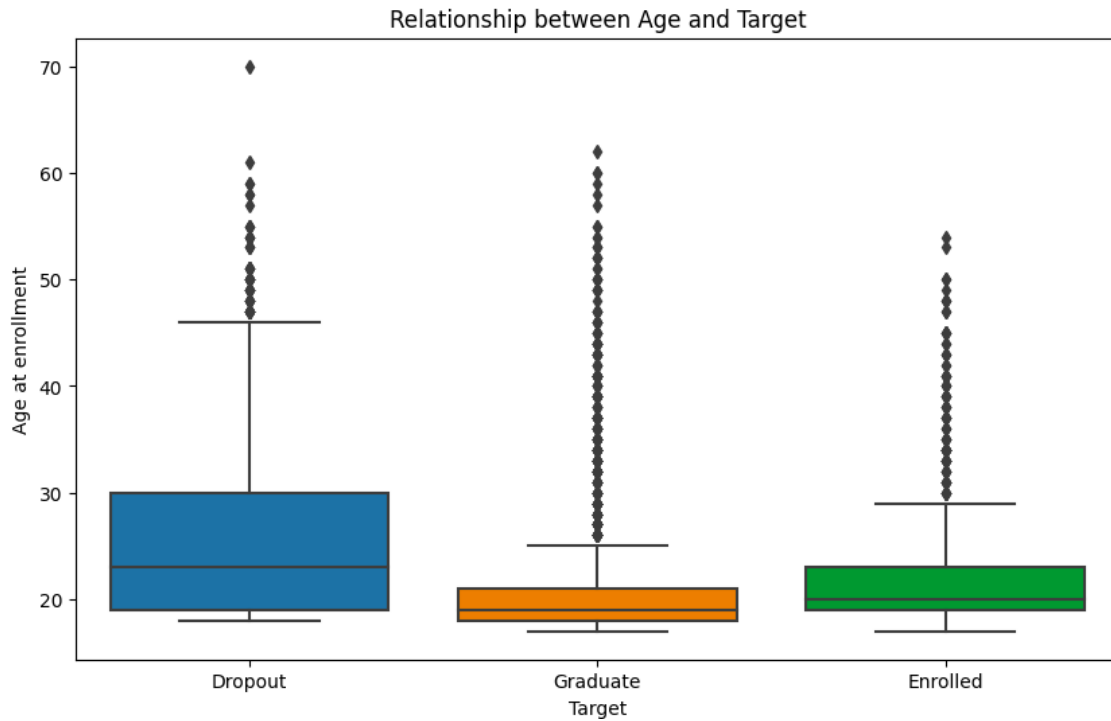
Creating a boxplot to showcase relationship between Age at enrollment and Target



```

[18] plt.figure(figsize=(10, 6))
sns.boxplot(x='Target', y='Age at enrollment', data=new_data)
plt.xlabel('Target')
plt.ylabel('Age at enrollment')
plt.title('Relationship between Age and Target')
plt.show()

```



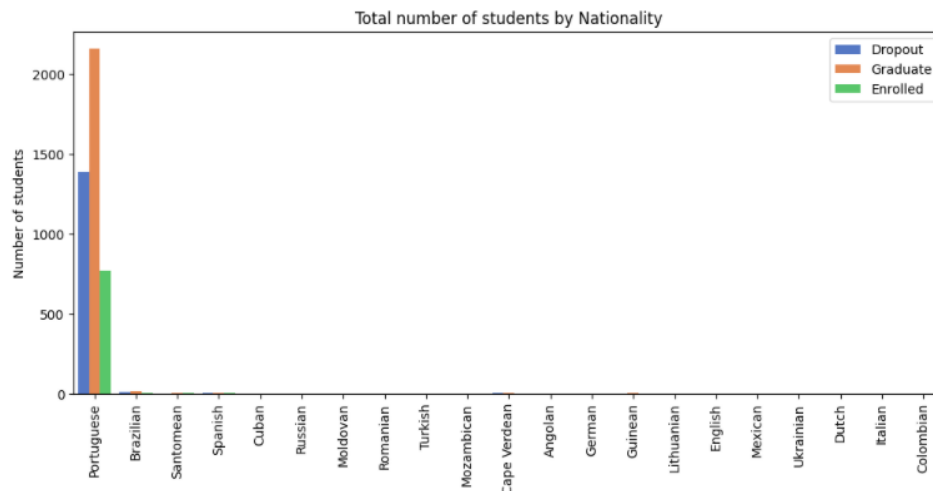
The highest graduation rate is the age group 17-22 with 72% of graduated students and 28% of dropped out students. There is no definitive distinction between other age groups.

## 5.8 Nationality

Creating a count chart showing the total number of dropout, graduate, and enrolled students based on Nationality.

```
[19] # Nationality
new_data[['Nationality']] = new_data[['Nationality']].replace({'Nationality': {1: 'Portuguese', 2: 'German', 6: 'Spanish',
11: 'Italian', 13: 'Dutch', 14: 'English',
17: 'Lithuanian', 21: 'Angolan', 22: 'Cape Verdean',
24: 'Guinean', 25: 'Mozambican', 26: 'Santomean',
32: 'Turkish', 41: 'Brazilian', 62: 'Romanian',
100: 'Moldovan', 101: 'Mexican', 103: 'Ukrainian',
105: 'Russian', 108: 'Cuban', 109: 'Colombian'}}})

fig, ax = plt.subplots(figsize=(12, 5))
order = new_data[new_data['Target'] == 'Enrolled']['Nationality'].value_counts()
ax = sns.countplot(data=new_data, x='Nationality', hue='Target', palette='muted', order=order.index)
ax.set(xlabel=None, ylabel='Number of students', title='Total number of students by Nationality')
plt.xticks(rotation=90)
ax.legend_.set_title(None)
plt.show()
```



The vast majority of students are Portuguese ( 4,314 , almost 98% of the whole dataset.). Hence , We can say that the variable “Nationality” is highly imbalanced.

Other than this, There is no pattern to be seen in the above plots.

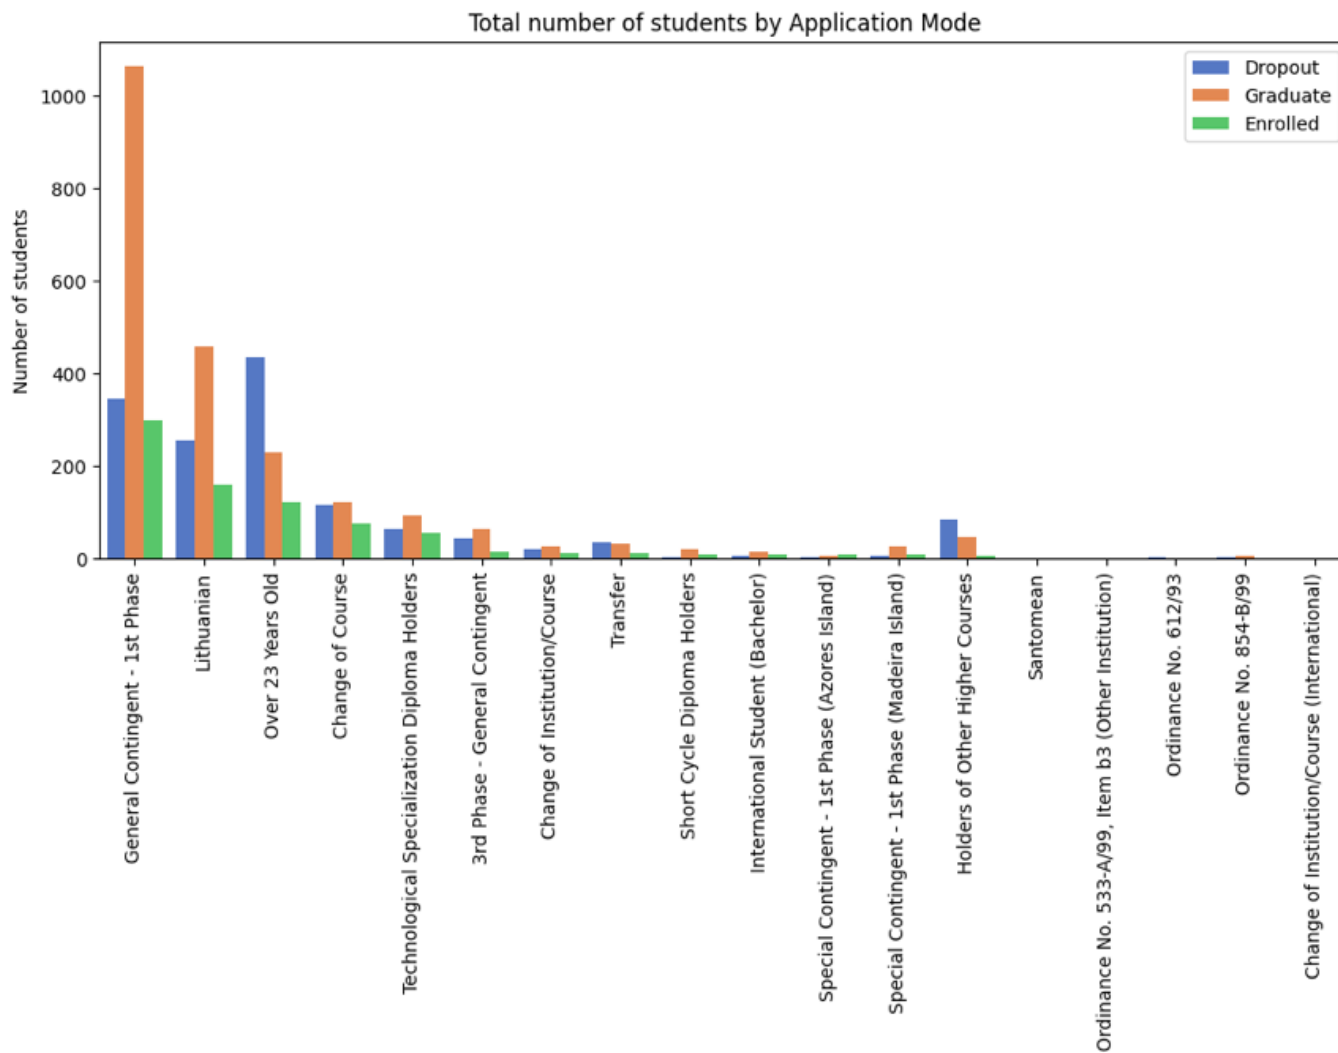
## 5.9 Application Mode

Creating a count plot that shows the proportion of dropout, and enrolled students in different application mode groups.

```
[20] # Application Mode

new_data['Application mode'] = new_data['Application mode'].replace({
    1: 'General Contingent - 1st Phase',
    2: 'Ordinance No. 612/93',
    5: 'Special Contingent - 1st Phase (Azores Island)',
    7: 'Holders of Other Higher Courses',
    10: 'Ordinance No. 854-B/99',
    15: 'International Student (Bachelor)',
    16: 'Special Contingent - 1st Phase (Madeira Island)',
    17: '2nd Phase - General Contingent',
    18: '3rd Phase - General Contingent',
    26: 'Ordinance No. 533-A/99, Item b2) (Different Plan)',
    27: 'Ordinance No. 533-A/99, Item b3 (Other Institution)',
    39: 'Over 23 Years Old',
    42: 'Transfer',
    43: 'Change of Course',
    44: 'Technological Specialization Diploma Holders',
    51: 'Change of Institution/Course',
    53: 'Short Cycle Diploma Holders',
    57: 'Change of Institution/Course (International)',
    6: 'Spanish',
    11: 'Italian',
    13: 'Dutch',
    14: 'English',
    17: 'Lithuanian',
    21: 'Angolan',
    22: 'Cape Verdean',
    24: 'Guinean',
    25: 'Mozambican',
    26: 'Santomean',
    32: 'Turkish',
    41: 'Brazilian',
    62: 'Romanian',
    100: 'Moldovan',
    101: 'Mexican',
    103: 'Ukrainian',
    105: 'Russian',
    108: 'Cuban',
    109: 'Colombian'
})

fig, ax = plt.subplots(figsize=(12, 5))
order = new_data[new_data['Target'] == 'Enrolled']['Application mode'].value_counts()
ax = sns.countplot(data=new_data, x='Application mode', hue='Target', palette='muted', order=order.index)
ax.set(xlabel=None, ylabel='Number of students', title='Total number of students by Application Mode')
plt.xticks(rotation=90)
ax.legend_.set_title(None)
plt.show()
```



The majority of currently enrolled students have 1st phase - general contingent (1) application mode and the graduation rate of these students is highest.

## 5.10 Daytime/Evening Attendance

Creating a pie chart that depicts the proportions of total number of daytime, and total number of evening students who dropped out.

```
[21] # Daytime/evening attendance

# Filter data for 'Target' equal to 0 (dropped out)
dropout_data = new_data.loc[new_data['Target'] == "Dropout"]

# Count the occurrences of each marital status
attendance_counts = dropout_data['Daytime/evening attendance'].value_counts().reset_index()

# Rename the columns for clarity in the plot
attendance_counts.columns = ['Daytime/evening attendance', 'Count']

# Calculate percentages
attendance_counts['Percentage'] = (attendance_counts['Count'] / attendance_counts['Count'].sum()) * 100

# Plotting the pie chart using plotly.express
fig = px.pie(attendance_counts, names='Daytime/evening attendance', values='Percentage', title='Daytime/evening attendance Distribution for Dropout Students', color_discrete_
fig.update_traces(labels=['Daytime', 'Evening'], hole=0.4, textinfo='percent+label', pull=[0.1, 0.1, 0.1, 0.1])

fig.show()
```

/usr/local/lib/python3.10/dist-packages/numpy/core/numeric.py:2463: FutureWarning:  
elementwise comparison failed; returning scalar instead, but in the future will perform elementwise comparison

Daytime/evening attendance Distribution for Dropout Students



Vast majority of students who dropped out studied during the daytime.  
The variable Daytime/evening attendance is imbalanced:

Note: It is not recommended to use Daytime/evening attendance as a predictor variable.

## 5.11 Course

Creating a stacked Bar-Plot that depicts the proportions of dropped out and Graduated students in a particular Course.


```
✓ [22] # Course
0s
new_data['Course'] = new_data['Course'].replace({
    33: 'Biofuel Production Technologies',
    171: 'Animation and Multimedia Design',
    8014: 'Social Service (Evening Attendance)',
    9003: 'Agronomy',
    9070: 'Communication Design',
    9085: 'Veterinary Nursing',
    9119: 'Informatics Engineering',
    9130: 'Equinculture',
    9147: 'Management',
    9238: 'Social Service',
    9254: 'Tourism',
    9500: 'Nursing',
    9556: 'Oral Hygiene',
    9670: 'Advertising and Marketing Management',
    9773: 'Journalism and Communication',
    9853: 'Basic Education',
    9991: 'Management (Evening Attendance)'
})

order = [
    'Biofuel Production Technologies',
    'Animation and Multimedia Design',
    'Social Service (Evening Attendance)',
    'Agronomy',
    'Communication Design',
    'Veterinary Nursing',
    'Informatics Engineering',
    'Equinculture',
    'Management',
    'Social Service',
    'Tourism',
    'Nursing',
    'Oral Hygiene',
    'Advertising and Marketing Management',
    'Journalism and Communication',
    'Basic Education',
    'Management (Evening Attendance)'
]

# Create a stacked bar plot.
filtered_data = pd.crosstab(index=new_data['Course'], columns=new_data[new_data['Target'] != 'Enrolled']['Target'])
data_prop = pd.crosstab(index=new_data['Course'], columns=new_data[new_data['Target'] != 'Enrolled']['Target'], normalize='index')
ax = data_prop.loc[order[:-1]].plot(kind='barh', stacked=True, figsize=(7, 6))
ax.set(xlabel='Proportion', ylabel='Course')
ax.spines(['right', 'top']).set_visible(False)
ax.legend_.set_title(None)
ax.set_title('Percentage of dropped out and graduated students')
plt.xticks(rotation=0)

for n, x in enumerate([*filtered_data.loc[order[:-1]].index.values]):
    for (proportion, count, y_loc) in zip(data_prop.loc[x], filtered_data.loc[x], data_prop.loc[x].cumsum()):
        plt.text(x=(y_loc - proportion) + (proportion / 5), y=n - 0.2, s=f'{count} ~ ({np.round(proportion * 100, 1)}%)')

plt.show()
```



The majority of students who successfully graduated were studying Nursing.

The highest dropout rates are in Informatics Engineering and Biofuel Production Technologies courses with 86.8% and 88.9% respectively.

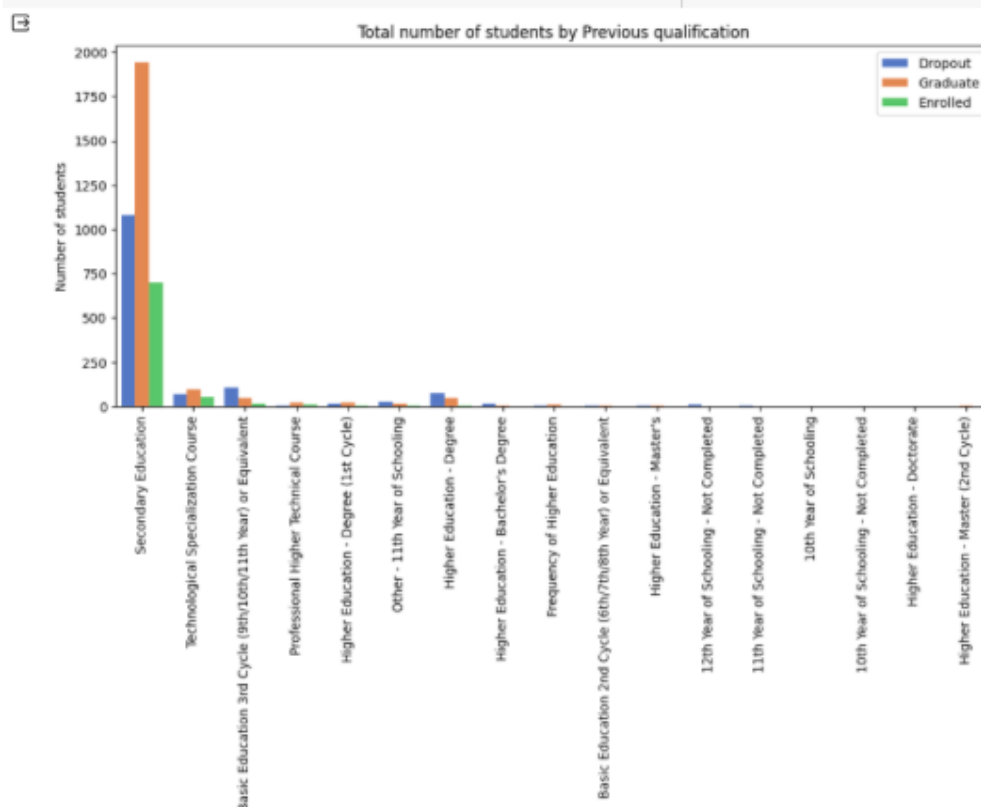
In general, there is a pattern and the percentage of dropout and graduate students in different courses does not stay the same and we can use Course as a Predictor for our model.

## 5.12 Previous Qualifications

Creating a Count plot showing relationship between Graduate, Dropout and Enrolled Students with Previous qualification.

```
# Previous Quaifications
1s
new_data['Previous qualification'] = new_data['Previous qualification'].replace({
    1: 'Secondary Education',
    2: 'Higher Education - Bachelor's Degree',
    3: 'Higher Education - Degree',
    4: 'Higher Education - Master's',
    5: 'Higher Education - Doctorate',
    6: 'Frequency of Higher Education',
    9: '12th Year of Schooling - Not Completed',
    10: '11th Year of Schooling - Not Completed',
    12: 'Other - 11th Year of Schooling',
    14: '10th Year of Schooling',
    15: '10th Year of Schooling - Not Completed',
    19: 'Basic Education 3rd Cycle (9th/10th/11th Year) or Equivalent',
    38: 'Basic Education 2nd Cycle (6th/7th/8th Year) or Equivalent',
    39: 'Technological Specialization Course',
    40: 'Higher Education - Degree (1st Cycle)',
    42: 'Professional Higher Technical Course',
    43: 'Higher Education - Master (2nd Cycle)'
})

fig, ax = plt.subplots(figsize=(12, 5))
order = new_data[new_data['Target'] == 'Enrolled']['Previous qualification'].value_counts()
ax = sns.countplot(data=new_data, x='Previous qualification', hue='Target', palette='muted', order=order.index)
ax.set(xlabel=None, ylabel='Number of students', title='Total number of students by Previous qualification')
plt.xticks(rotation=90)
ax.legend_.set_title(None)
plt.show()
```



The vast majority of students have secondary education (1) as Previous Qualification. The variable Previous qualification is highly imbalanced. Moreover, There is no pattern to be seen in the above plots.



## For Continuous data

- 5.13.1 Curricular units 1st sem (credited)
- 5.13.2 Curricular units 1st sem (enrolled)
- 5.13.3 Curricular units 1st sem (evaluations)
- 5.13.4 Curricular units 1st sem (approved)
- 5.13.5 Curricular units 1st sem (grade)
- 5.13.6 Curricular units 1st sem (without evaluations)
- 5.13.7 Curricular units 2nd sem (credited)
- 5.13.8 Curricular units 2nd sem (enrolled)
- 5.13.9 Curricular units 2nd sem (evaluations)
- 5.13.10 Curricular units 2nd sem (approved)
- 5.13.11 Curricular units 2nd sem (grade)
- 5.13.12 Curricular units 2nd sem (without evaluations)
- 5.13.13 Unemployment rate
- 5.13.14 Inflation rate
- 5.13.15 GDP

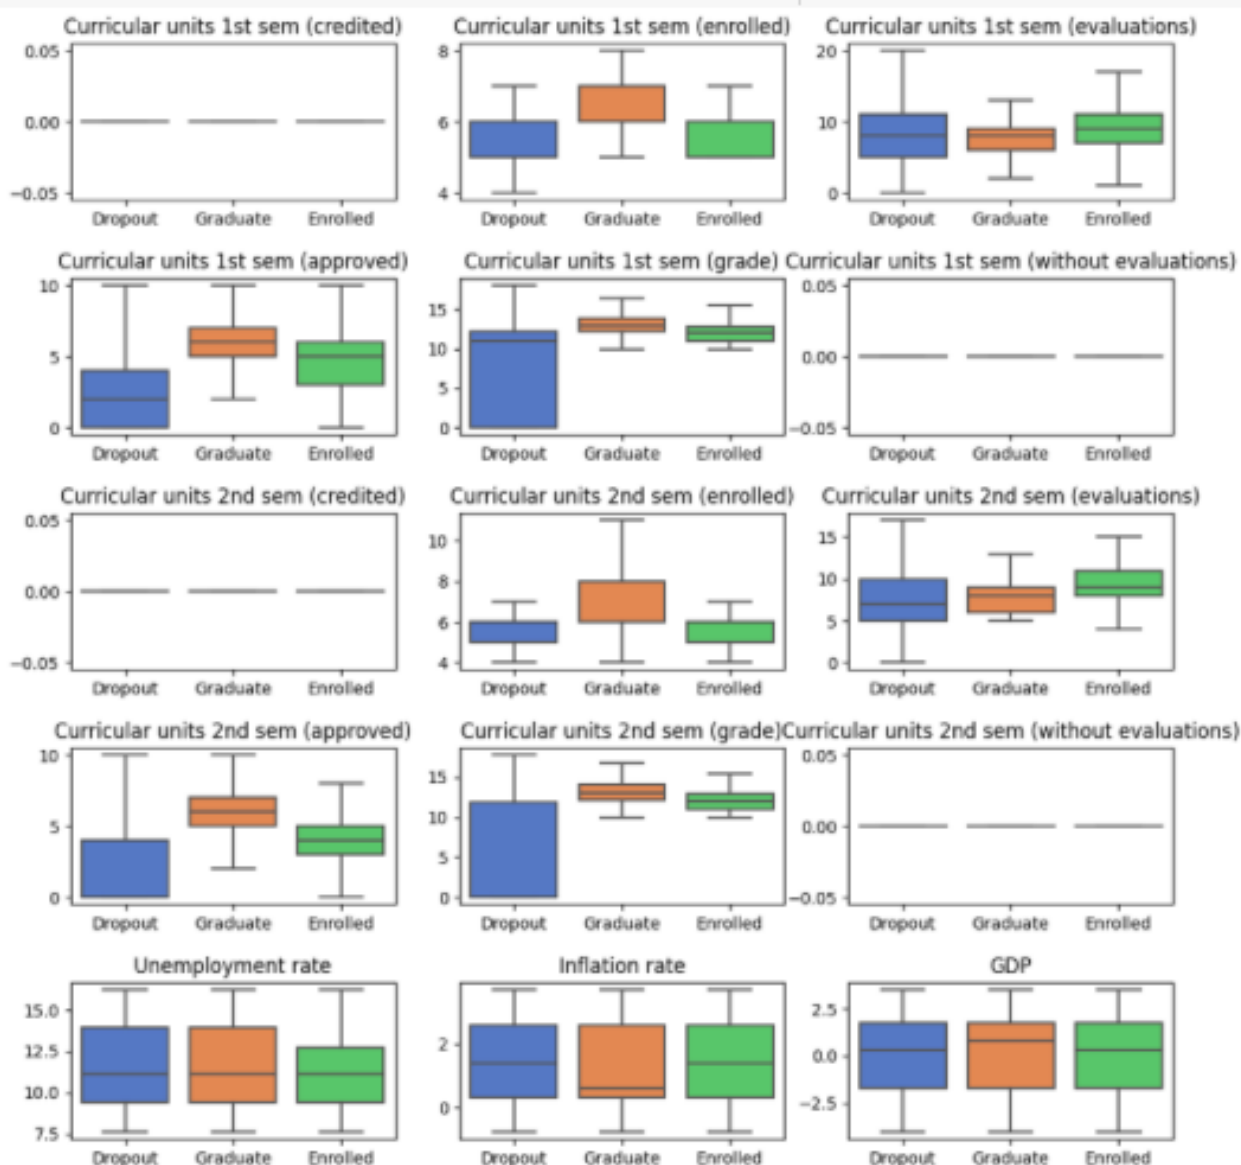
Creating a box plot for all the above mentioned continuous data.

```

✓ [24] # Continuous Variables
1s
fig, axs = plt.subplots(5, 3, figsize=(12, 12))
plt.subplots_adjust(hspace=0.5)
cont_cols = ['Curricular units 1st sem (credited)',
             'Curricular units 1st sem (enrolled)',
             'Curricular units 1st sem (evaluations)',
             'Curricular units 1st sem (approved)',
             'Curricular units 1st sem (grade)',
             'Curricular units 1st sem (without evaluations)',
             'Curricular units 2nd sem (credited)',
             'Curricular units 2nd sem (enrolled)',
             'Curricular units 2nd sem (evaluations)',
             'Curricular units 2nd sem (approved)',
             'Curricular units 2nd sem (grade)',
             'Curricular units 2nd sem (without evaluations)',
             'Unemployment rate',
             'Inflation rate',
             'GDP']

ind = 0
for i in range(5):
    for j in range(3):
        sns.boxplot(new_data, x='Target', y=cont_cols[ind], showfliers=False, palette='muted', ax=axs[i, j])
        axs[i, j].set(xlabel=None, ylabel=None, title=cont_cols[ind])
        ind += 1

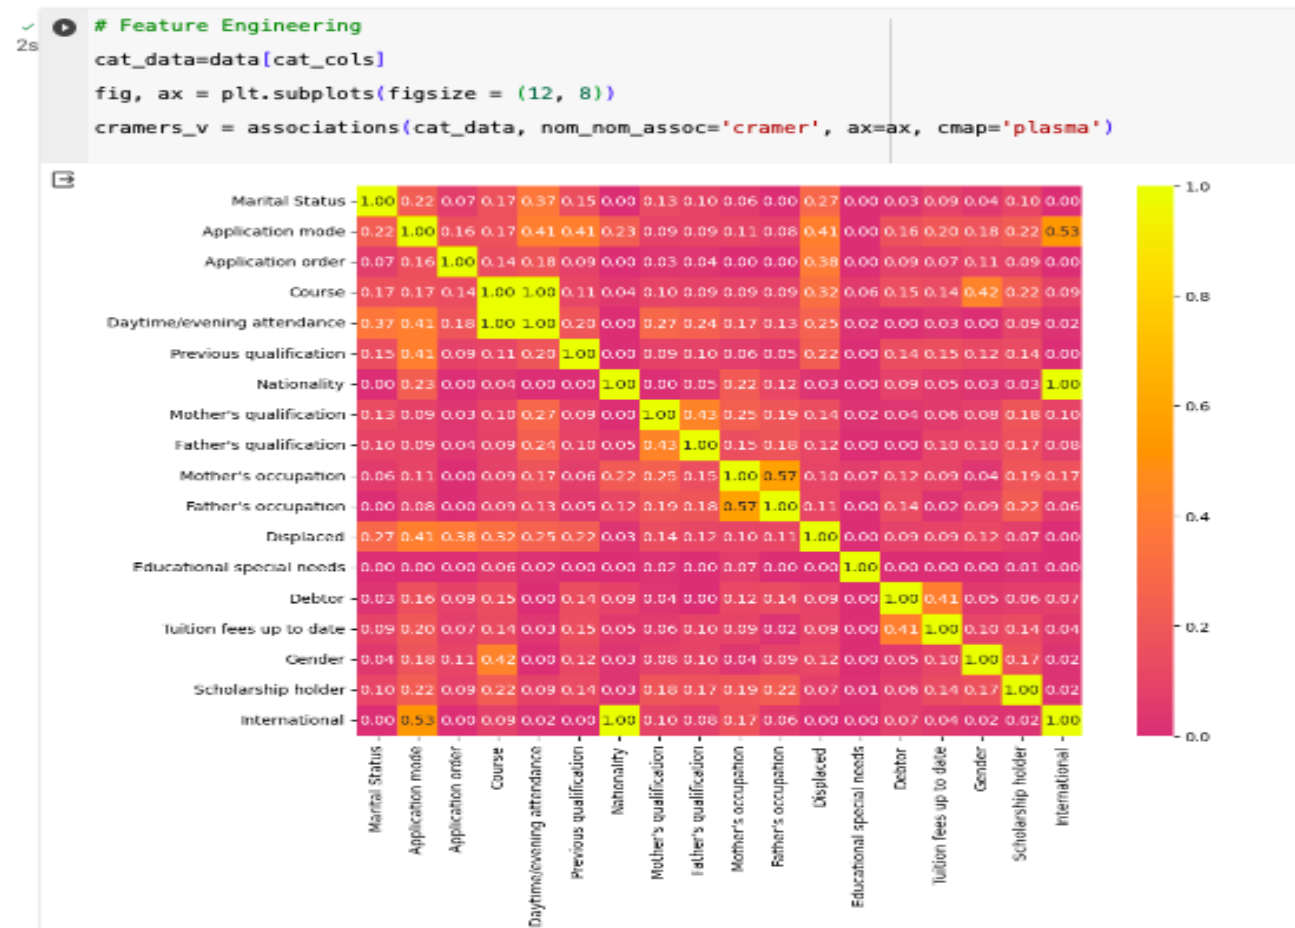
```



## 6. Feature Engineering

### For Categorical data

Creating a correlation matrix for the categorical data.



We have some features which have solid association with each other, these features are redundant and will not provide the model new information.

Mother's occupation and Father's occupation have good association with each other (0.57). Keep only Mother's occupation as it has higher association with the Target.

Mother's qualification and Father's qualification have good association with each other (0.43). Keep only Mother's qualification as it has higher association with the Target.

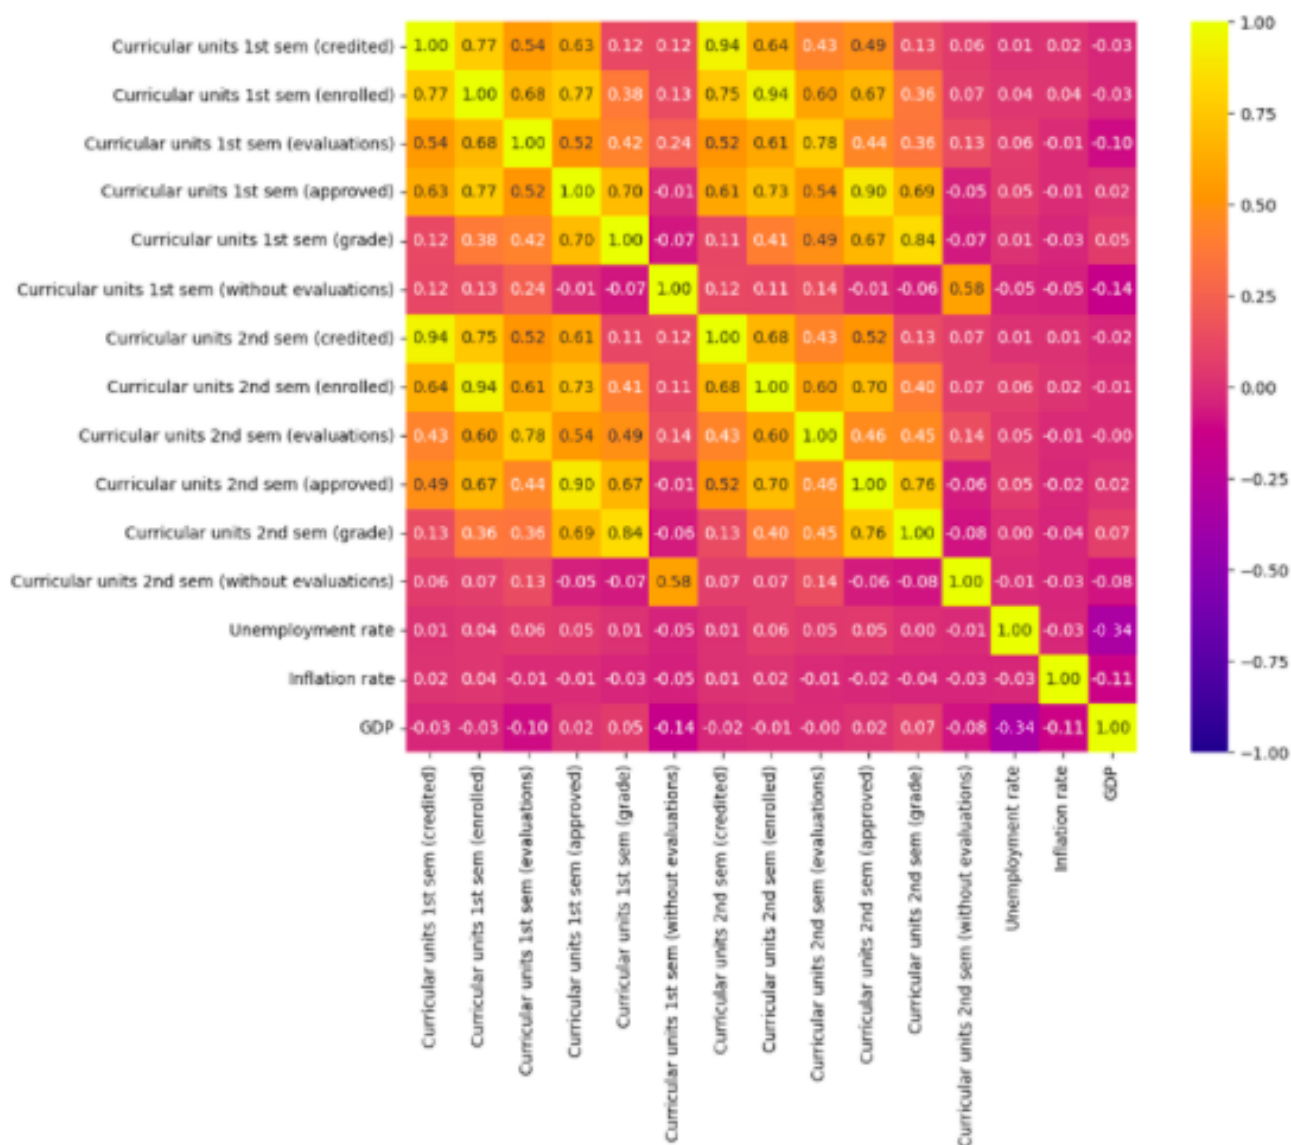
Debtor and Tuition fees up to date have good association with each other (0.41). Keep only Tuition fees up to date as it has higher association with the Target.

Displaced and Application mode have good association with each other (0.41). Keep only Application mode as it has higher association with the Target.

## For Continuous data

Creating a correlation matrix for the continuous data.

```
✓ [26] cont_data=data[cont_cols]
1s fig, ax = plt.subplots(figsize = (12, 8))
cramers_v = associations(cont_data, nom_nom_assoc='cramer', ax=ax, cmap='plasma')
```



According to the correlation matrix we will select following Categorical features: Application mode, Course, Previous qualification, Mother's qualification, Tuition fees up to date, Mother's occupation, Gender, Scholarship holder,

According to the correlation matrix we will select following Continuous features: Age at enrollment, Curricular units 1st sem (approved), Curricular units 2nd sem (approved).

```
selected_cols = ['Application mode', 'Course', 'Previous qualification', "Mother's qualification", 'Tuition fees up to date',  
                "Mother's occupation", 'Gender', 'Scholarship holder', 'Age at enrollment', 'Curricular units 1st sem (approved)',  
                'Curricular units 2nd sem (approved)', 'Target']  
  
# Keep only relevant columns.  
new_data=new_data[selected_cols]  
  
# Remove enrolled students.  
new_data = new_data[new_data['Target'] != 'Enrolled']  
  
# Convert into numerical data type.  
  
cols = ['Tuition fees up to date', 'Gender', 'Scholarship holder', 'Target']  
new_data[cols] = new_data[cols].astype('int32')
```

Now, The dataset is ready for model building.

## 7. Model Building

We will be apply the following ML classification algorithms :

- Logistic regression
- Decision Tree Classifier
- Random Forest Classifier
- Naive Bayes Classifier
- Support Vector Machine

The motivation to apply all these algorithms was that we wanted to compare their accuracy results to see which algorithm works better on our dataset.

Splitting the dataset into training and testing set with (80% , 20%) ratio respectively, and evaluating the models based on their accuracy, precision, recall, F1 score



```
# Model Building

results = pd.DataFrame(columns=['Algorithm', 'Accuracy', 'Precision', 'Recall', 'F1 Score'])

# Predicting variable.
y = new_data['Target']

# Predictor features.
X = new_data.copy()
X = X.drop('Target', axis = 1)

# Create training and test sets, 80% and 20% respectively.
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=42)
```

For each case, we've visualized the Confusion Matrix along with it. We've also displayed the accuracy percentage for each case too.

In the end we have displayed the combined ROC graph of these ML classification algorithms.

## 8. Models

### 8.1 Naive Bayes Classifier

```
# Naive Bayes

# Normalize data.
X_scaled_train = StandardScaler().fit_transform(X_train)
X_scaled_test = StandardScaler().fit_transform(X_test)
# Fit the model.
gnb = GaussianNB()

gnb.fit(X_scaled_train, y_train)
# Get the predictions on test data.
y_preds = gnb.predict(X_scaled_test)

print_results('Naive Bayes', y_test, y_preds)

num_iterations = 20
accuracy_scores = []

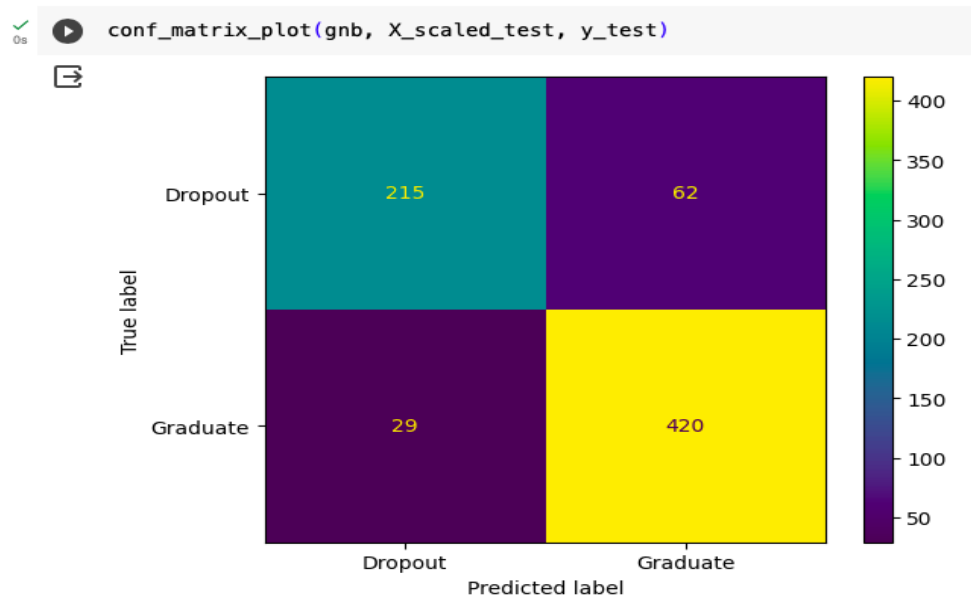
y_preds_prob_gnb = gnb.predict_proba(X_test)[: , 1]

for i in range(num_iterations):
    scores = cross_val_score(gnb, X_scaled_train, y_train, cv=20, scoring='accuracy')
    accuracy_scores.extend(scores)

# Print the average accuracy
average_accuracy = sum(accuracy_scores) / len(accuracy_scores)
print(f"Average Accuracy over {num_iterations} iterations:", average_accuracy)
```



```
Naive Bayes
Accuracy: 0.875
Precision: 0.871
Recall: 0.935
F1 Score: 0.902
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning:
X has feature names, but GaussianNB was fitted without feature names
```



## Naive Bayes

- Accuracy: 0.875
- Precision: 0.871
- Recall: 0.935
- F1 Score: 0.902
- **Average Accuracy over 20 iterations: 0.8636230514879547**

Why we are using this:

- Naive Bayes assumes independence between features, which can be a strength if this assumption approximately holds.
- Suitable for high-dimensional datasets.
- Naive Bayes might be effective if the features are conditionally independent given the class (dropout or academic success).
- It's computationally efficient and can work well with large datasets.



## 8.2 Logistic Regression

```
[33] # Logistic Regression

X_scaled_train = StandardScaler().fit_transform(X_train)
X_scaled_test = StandardScaler().fit_transform(X_test)
lr = LogisticRegression()
lr.fit(X_scaled_train, y_train)
y_preds = lr.predict(X_scaled_test)

print_results('Logistic Regression', y_test, y_preds)

num_iterations = 20
accuracy_scores = []

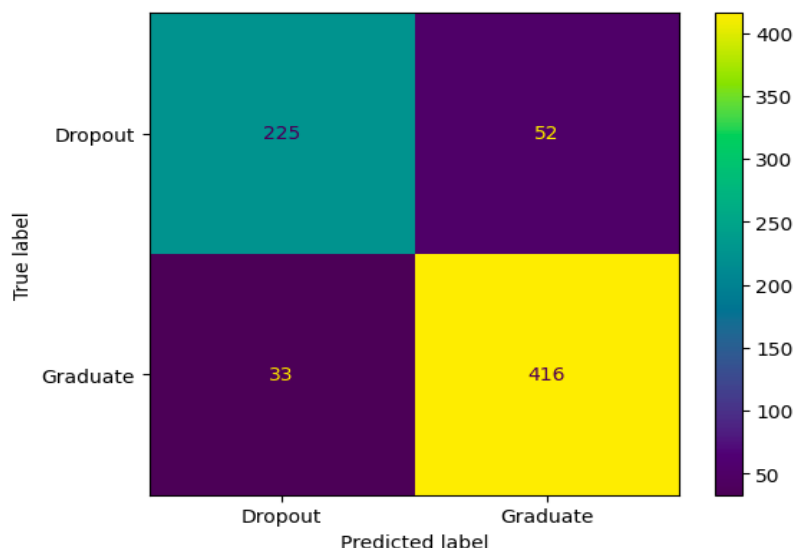
for i in range(num_iterations):
    scores = cross_val_score(lr, X_scaled_train, y_train, cv=20, scoring='accuracy')
    accuracy_scores.extend(scores)

y_preds_prob_lr = lr.predict_proba(X_scaled_test)[: , 1]

# Print the average accuracy
average_accuracy = sum(accuracy_scores) / len(accuracy_scores)
print(f"Average Accuracy over {num_iterations} iterations:", average_accuracy)
```

```
Logistic Regression
Accuracy: 0.883
Precision: 0.889
Recall: 0.927
F1 Score: 0.907
Average Accuracy over 20 iterations: 0.8894544166273
```

```
✓ [57] conf_matrix_plot(lr, X_scaled_test, y_test)
```



### Logistic Regression

- Accuracy: 0.883
- Precision: 0.889
- Recall: 0.927
- F1 Score: 0.907
- **Average Accuracy over 20 iterations: 0.8894544166273**

Why we are using this:

- Logistic Regression assumes a linear relationship between the independent variables and the log-odds of the dependent variable.
- Suitable when the classes are linearly separable.
- The linear relationship assumption might be a good fit for your dataset.
- Logistic Regression is less prone to overfitting, which could be beneficial if your dataset is not very large.

## 8.3 Decision Tree



# Decision Tree

```
decision_tree = DecisionTreeClassifier(random_state=0)
decision_tree.fit(X_train, y_train)
y_preds = decision_tree.predict(X_test)

print_results('Decision Tree', y_test, y_preds)

num_iterations = 20
accuracy_scores = []

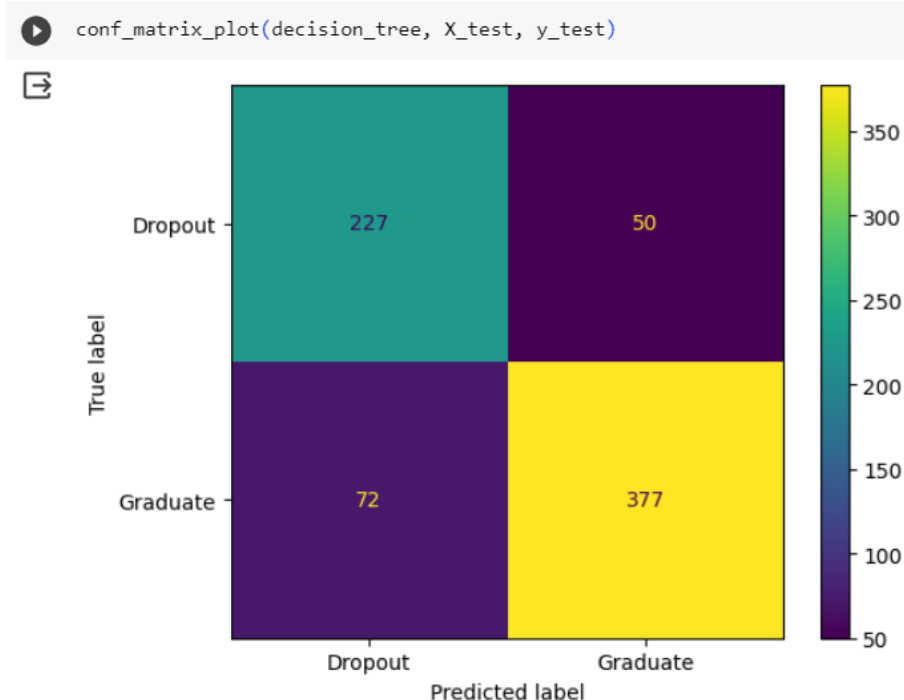
for i in range(num_iterations):
    scores = cross_val_score(decision_tree, X_scaled_train, y_train, cv=20, scoring='accuracy')
    accuracy_scores.extend(scores)

y_preds_prob_dt = decision_tree.predict_proba(X_test)[: , 1]

# Print the average accuracy
average_accuracy = sum(accuracy_scores) / len(accuracy_scores)
print(f"Average Accuracy over {num_iterations} iterations:", average_accuracy)
```



```
Decision Tree
Accuracy: 0.832
Precision: 0.883
Recall: 0.840
F1 Score: 0.861
Average Accuracy over 20 iterations: 0.8401771374586703
```




### Decision Tree

- Accuracy: 0.832
- Precision: 0.883
- Recall: 0.840
- F1 Score: 0.861
- **Average Accuracy over 20 iterations: 0.8401771374586703**

### Why we are using this:

- Decision Trees can capture non linear relationships in the data.
- They are easy to interpret.
- Decision Trees might struggle if the dataset has complex relationships that are hard to capture with a single split at each node.
- Prone to overfitting, especially if the tree is deep and the dataset is not large enough.

## 8.4 Random Forest

```
 # Random Forest

rf = RandomForestClassifier(random_state=0)
rf.fit(X_train, y_train)
y_preds = rf.predict(X_test)


print_results('Random forest', y_test, y_preds)

num_iterations = 20
accuracy_scores = []

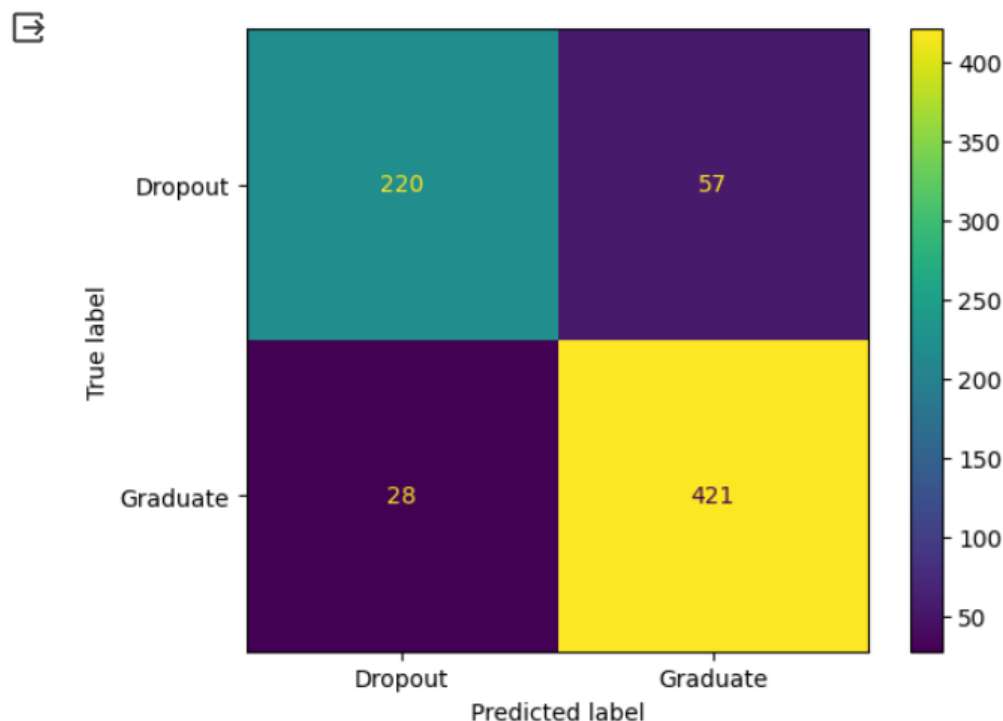
y_preds_prob_rf = rf.predict_proba(X_test)[: , 1]

for i in range(num_iterations):
    scores = cross_val_score(rf, X_scaled_train, y_train, cv=20, scoring='accuracy')
    accuracy_scores.extend(scores)

# Print the average accuracy
average_accuracy = sum(accuracy_scores) / len(accuracy_scores)
print(f"Average Accuracy over {num_iterations} iterations:", average_accuracy)
```

```
 Random forest
Accuracy: 0.883
Precision: 0.881
Recall: 0.938
F1 Score: 0.908
Average Accuracy over 20 iterations: 0.8880538497874321
```

```
conf_matrix_plot(rf, X_test, y_test)
```



### Random forest

- Accuracy: 0.883
- Precision: 0.881
- Recall: 0.938
- F1 Score: 0.908
- **Average Accuracy over 20 iterations: 0.8880538497874321**

### Why we are using this:

- Random Forest is an ensemble method that can handle non-linearity and complex relationships in the data.
- Effective in dealing with irrelevant features and outliers.
- Random Forest might capture complex interactions between attributes in predicting student outcomes.
- It can handle a mix of numerical and categorical features well.

## 8.5 Support Vector Machine (SVM)



# SVM

```
X_scaled_train = StandardScaler().fit_transform(X_train)
X_scaled_test = StandardScaler().fit_transform(X_test)
svm = SVC(probability=True)
svm.fit(X_scaled_train, y_train)
y_preds = svm.predict(X_scaled_test)

print_results('SVM', y_test, y_preds)

num_iterations = 20
accuracy_scores = []

y_preds_prob_svm = svm.predict_proba(X_scaled_test)[: , 1]


for i in range(num_iterations):
    scores = cross_val_score(svm, X_scaled_train, y_train, cv=20, scoring='accuracy')
    accuracy_scores.extend(scores)

# Print the average accuracy
average_accuracy = sum(accuracy_scores) / len(accuracy_scores)
print(f"Average Accuracy over {num_iterations} iterations:", average_accuracy)
```



```
SVM
Accuracy: 0.888
Precision: 0.872
Recall: 0.960
F1 Score: 0.914
Average Accuracy over 20 iterations: 0.898398677373639
```

```
conf_matrix_plot(svm, X_scaled_test, y_test)
```



## SVM

- Accuracy: 0.888
- Precision: 0.872
- Recall: 0.960
- F1 Score: 0.914
- **Average Accuracy over 20 iterations: 0.898398677373639**

Why we are using this:

- SVM works well when there is a clear margin of separation between classes.
- Effective in high-dimensional spaces, which is beneficial if your dataset has many features.
- The dataset might have a clear separation between students who drop out and those who succeed academically.
- SVM is robust to outliers, and if there are outliers in the data, SVM can handle them effectively.



## 9. Model Evaluation and Conclusion

### 9.1 Result matrix for different ML classification algorithms

```
✓ [80] # Compare models.
0s results_unique=results.drop_duplicates()
    results_unique.sort_values(by=['Accuracy'], ascending=False)
```

	Algorithm	Accuracy	Precision	Recall	F1 Score
4	SVM	0.888430	0.872470	0.959911	0.914104
1	Logistic Regression	0.882920	0.888889	0.926503	0.907306
3	Random forest	0.882920	0.880753	0.937639	0.908306
0	Naive Bayes	0.874656	0.871369	0.935412	0.902256
2	Decision Tree	0.831956	0.882904	0.839644	0.860731

From the above table , it can be clearly observed that SVM, Logistic Regression and Random Forest are the models with highest Accuracy, Precision, Recall and F1 Score. But only the above parameters might not be sufficient to arrive at our decision for best classifier. Let's explore another approach.

## 9.2 ROC Graph

```


y_preds_probs = {'Naive Bayes': y_preds_prob_gnb, 'Logistic regression': y_preds_prob_lr, 'SVM': y_preds_prob_svm,
                  'Decision Tree': y_preds_prob_dt, 'Random forest': y_preds_prob_rf}

aucs = {}
for i in y_preds_probs.keys():
    aucs[i] = metrics.roc_auc_score(y_test, y_preds_probs[i])

for i in dict(sorted(aucs.items(), key=operator.itemgetter(1), reverse=True)).keys():
    fpr, tpr, thresholds = metrics.roc_curve(y_test, y_preds_probs[i])
    auc = aucs[i]
    plt.plot(fpr, tpr, linestyle='dashdot', label=f'{i}, AUC = {auc:.3f}')

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.legend(loc='best')
plt.show()

```



**9.3 Conclusion:** Here, Random Forest has the highest AUC (Area Under Curve) among all models. Hence, now we can confidently say that Random Forest is the best Classifier for our Dataset.

## 10. References:

1. M.V.Martins, D. Tolledo, J. Machado, L. M.T. Baptista, V.Realinho. (2021) "Early prediction of student's performance in higher education: a case study" Trends and Applications in Information Systems and Technologies, vol.1, in Advances in Intelligent Systems and Computing series. Springer. DOI: 10.1007/978-3-030-72657-7\_16 , Source of Dataset: [Predict students' dropout and academic success](#)
2. Official Documentation of [Scikit-Learn Documentation](#), [Matplotlib Documentation](#), [Pandas Documentation](#), [Seaborn Documentation](#), [NumPy Documentation](#), [Operator Module Documentation](#), [Dython Documentation](#)
3. Class Notes of IDS Course.
4. [Machine Learning by Andrew Ng on Coursera](#), [Towards Data Science on Medium](#)